



Attention Divergence from BiMGA (Full), ordered by increasing divergence

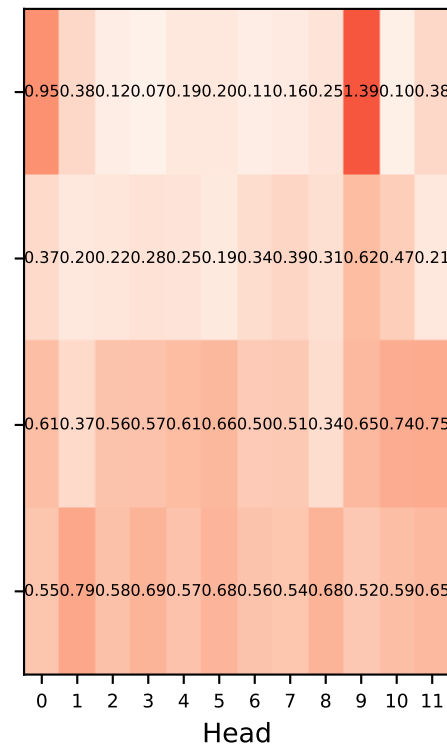
BiMGA Uniform
MRR=0.313 | Mean KL=0.086



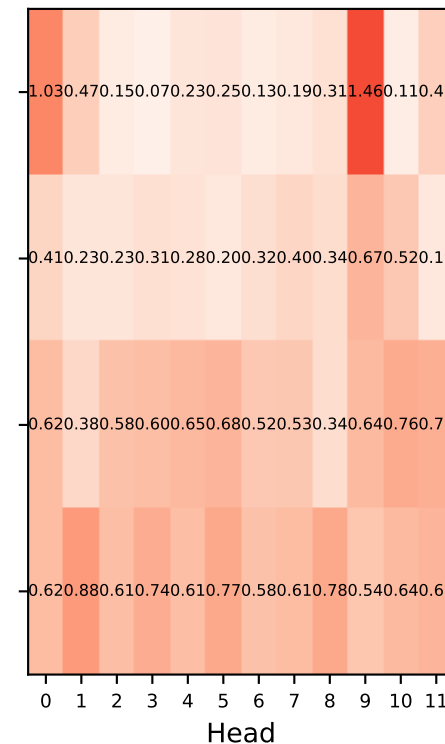
Embed Distill
MRR=0.303 | Mean KL=0.210



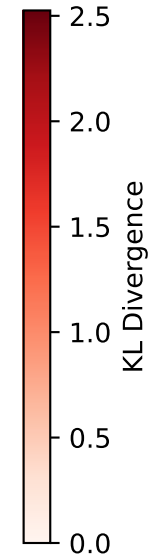
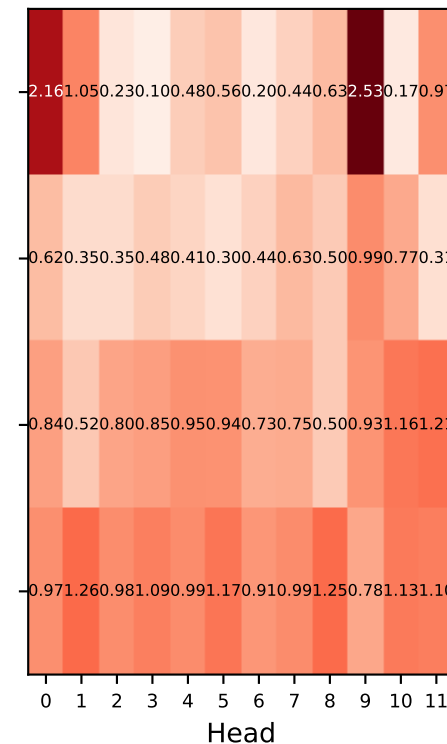
Score Distill
MRR=0.301 | Mean KL=0.467



Hard Neg Pairwise
MRR=0.302 | Mean KL=0.500



Control (Supervised)
MRR=0.205 | Mean KL=0.801



Models with explicit alignment objectives (BiMGA Uniform, Embed Distill) develop attention patterns closest to BiMGA. Output-matching methods (Score Distill, Hard Neg Pairwise) and supervised training diverge most, especially in layers 0 and 2-3.